

WENZHE ZHU

✉ wzhzhu@ustc.edu.cn · 📞 (+86) 18855726088 · 🌐 Homepage

Availability: Fall 2027 for full-time postdoctoral appointment; open to early remote collaboration or short-term research visits.

RESEARCH PROFILE

I am a postdoctoral fellow at the University of Science and Technology of China. My research focuses on building scalable and resource-efficient storage, memory, and data-serving systems for data-intensive computing and emerging AI workloads. My work spans distributed storage systems, key-value stores, disaggregated memory, and checkpointing systems for large language model training. My research has appeared in top-tier systems and data management venues, including *SIGMETRICS*, *OSDI*, *FAST*, *ICDE*, *ICS*, *ICPP*, and *SRDS*.

EDUCATION

University of Science and Technology of China, Hefei, China 2018 – 2025

Ph.D. in Computer Science

Thesis: Performance and Overhead Optimization of Log-Structured Cloud Storage Systems

Advisors: Prof. Yinlong Xu and Prof. Yongkun Li

Hunan University, Changsha, China 2014 – 2018

B.Eng. in Software Engineering

RESEARCH EXPERIENCE & PROJECTS

Distributed and Cloud Storage Systems

- **Traffic scheduling for cloud block storage** 2020 – 2022
Studied traffic imbalance and tail-latency amplification in large-scale cloud block storage systems under bursty workloads. Developed differentiated forwarding mechanisms to balance traffic across storage nodes and reduce tail latency. Published in *ACM SIGMETRICS 2023*.
- **Data passing for stateful serverless analytics** 2022 – 2023
Identified intermediate data movement as a major performance and cost bottleneck in I/O-intensive stateful serverless analytics. Designed a high-performance, cost-efficient data-passing mechanism that reduces data-movement overhead and improves end-to-end execution efficiency. Published in *FAST 2024*.
- **Load balancing for distributed key-value stores** 2023 – 2025
Characterized cross-machine load imbalance in distributed key-value stores caused by heterogeneous CPU and disk-I/O demands under mixed workloads. Designed a cooperative CPU-I/O scheduling mechanism that coordinates compute and storage resources across nodes, improving load balance, throughput, and tail latency. Published in *ICDE 2026*.
- **Efficient erasure-coded key-value storage** 2025 – 2026
Led the study of update and repair overheads in erasure-coded key-value stores under dynamic workloads. Developed dynamic decoupling mechanisms that bypass update and repair penalties while preserving storage efficiency and fault tolerance. Accepted to appear in *SRDS 2026*.

Storage Optimizations for Data-Intensive Applications

- **Incremental checkpointing for LLM training** 2024 – 2025
Led the characterization of storage and I/O bottlenecks introduced by frequent checkpointing in large language model training. Designed and implemented a fast, space-efficient incremental checkpointing mechanism that reduces checkpoint storage overhead and checkpointing time. Published in *ICPP 2025*.

- **Resource-efficient key-value engine optimization** 2025 – Present
Led the study of cache inefficiency and concurrency overheads in LSM-tree-based key-value engines under constrained memory resources. Developed cooperative metadata–data cache management and hierarchical caching mechanisms that improve cache utilization and read performance. Accepted to appear in *ICS 2026*; one related manuscript is ready for submission.
- **Resource-efficient vector indexing for AI retrieval systems** 2025 – Present
Leading the design of memory- and scalability-efficient vector indexing mechanisms for AI-driven retrieval workloads, including retrieval-augmented generation and semantic search. Exploring hybrid clustering–graph index designs that reduce index memory overhead while preserving high recall, together with topology-aware shard partitioning mechanisms that reduce scatter-gather query overhead in distributed vector databases. Ongoing work.

Disaggregated Memory Systems

- **Fine-grained disaggregated memory management** 2024 – 2025
Studied the trade-off between allocation overhead and memory waste in fine-grained disaggregated memory systems. Designed a memory management framework that reduces allocation overhead while preserving fine-grained memory efficiency. Published in *OSDI 2025*.
- **Elastic and low-latency remote swap for disaggregated memory** 2025 – Present
Investigated RDMA-based remote swap for transparent memory expansion in disaggregated memory architectures. Designed elastic fine-grained memory slot allocation to reduce memory stranding caused by static peak-demand provisioning, and redesigned swap metadata management with reference-count arrays and hierarchical queues to achieve $O(1)$ swap entry allocation. Manuscript is ready for submission.

PROFESSIONAL EXPERIENCE

University of Science and Technology of China, Hefei, China Aug. 2025 – Jul. 2027 (*expected*)

Postdoctoral Fellow

Working with Prof. Yongkun Li.

Alibaba Cloud, Beijing, China

Nov. 2019 – Jun. 2021

Research Intern

Research topic: Low-latency and fairness-aware traffic scheduling in distributed block storage systems.

PUBLICATIONS

* indicates lead author / project lead and corresponding author.

1. Xiyuan Liang, **Wenzhe Zhu***, Chang Feng, Weiqiang Yu, Xinran Zhang, and Weifa Liang. “XORbit: Bypassing Update and Repair Penalties in Erasure-Coded KV Stores via Dynamic Decoupling.” To appear in *International Symposium on Reliable Distributed Systems (SRDS 2026)*.
2. Haoting Tang, **Wenzhe Zhu***, Qingyang Zhang, Jiakun Zhang, Junlin Jiang, Zaigui Zhang, Haijun Zhang, Yongkun Li, and Yinlong Xu. “CoCache: Accelerating Reads in KV Stores via Cooperative Metadata and Data Cache Management.” To appear in *ACM International Conference on Supercomputing (ICS 2026)*.
3. Jiakun Zhang, Patrick P. C. Lee, **Wenzhe Zhu**, Yongkun Li, Shuyi Zhang, and Yinlong Xu. “Cooperative Scheduling Between CPU and Disk I/O for Load Balancing in Distributed Key-Value Stores.” *IEEE International Conference on Data Engineering (ICDE 2026)*.
4. Zhiqiang Wang, **Wenzhe Zhu***, Zaigui Zhang, Chaomei Yan, Fan Guo, Yongkun Li, and Yinlong Xu. “Amber: Towards Fast and Space-Efficient Incremental Checkpointing in Large Language Model Training.” *International Conference on Parallel Processing (ICPP 2025)*.
5. Shuai Lin, Rui Wang, Zaigui Zhang, Long Deng, **Wenzhe Zhu**, Yongkun Li, and Yinlong Xu. “PTWalker: Cache-Efficient Random Walks via Alternating Dual-Subgraph Walker Updating.” *International Conference on Parallel Processing (ICPP 2025)*.

6. Xiaoyang Wang, Yongkun Li, Kan Wu, **Wenzhe Zhu**, Yuqi Li, and Yinlong Xu. “FineMem: Breaking the Allocation Overhead vs. Memory Waste Dilemma in Fine-Grained Disaggregated Memory Management.” *USENIX Symposium on Operating Systems Design and Implementation (OSDI 2025)*.
7. Tao Li, Yongkun Li, **Wenzhe Zhu**, Yinlong Xu, and John C. S. Lui. “MinFlow: High-Performance and Cost-Efficient Data Passing for I/O-Intensive Stateful Serverless Analytics.” *USENIX Conference on File and Storage Technologies (FAST 2024)*.
8. **Wenzhe Zhu**, Yongkun Li, Erci Xu, Fei Li, Yinlong Xu, and John C. S. Lui. “DiffForward: On Balancing Forwarding Traffic for Modern Cloud Block Services via Differentiated Forwarding.” *ACM SIGMETRICS International Conference on Measurement and Modeling of Computer Systems (SIGMETRICS 2023)*.
9. Haoyu Mao, Yongkun Li, **Wenzhe Zhu**, Fei Li, and Yinlong Xu. “On Optimizing Traffic Imbalance in Large-Scale Block-Based Cloud Storage: Trace Analysis and Algorithm Design.” *IEEE International Conference on Parallel and Distributed Systems (ICPADS 2022)*.

FUTURE RESEARCH DIRECTIONS

Going forward, I aim to design storage-centric system support for LLM-based agents deployed across end devices and edge servers, improving service quality and resource efficiency under tight memory, storage, compute, and latency constraints. This direction connects my background in storage and memory systems with emerging edge AI workloads, where efficient data management is critical for supporting retrieval, reasoning, and memory operations close to users. In particular, I plan to explore the following directions:

- **Resource-efficient vector databases:** designing compact vector indexes to reduce memory footprint and improve scalability for resource-constrained retrieval workloads on edge servers and end devices.
- **Storage-aware agent serving:** coordinating retrieval, context construction, caching, and inference-side data access across end devices, edge servers, and remote servers to balance latency, resource efficiency, and answer quality in agent services.
- **Agent memory management:** building efficient memory and storage support for organizing, persisting, retrieving, and evicting long-term, personalized, and task-specific agent memory.

TECHNICAL SKILLS

- **Systems:** Distributed storage, key-value stores, vector indexing, disaggregated memory, RDMA, Linux kernel.
- **Programming:** C/C++, Python, Go, Shell scripting.